

AKASH AGALAVE

Pune, India

☎ +91-9359372603

✉ akash.agalave.work@gmail.com

🌐 [linkedin.com/in/akash-agalave](https://www.linkedin.com/in/akash-agalave)

🐙 github.com/akashagalave

EDUCATION

Dr. D.Y. Patil University

Dec 2021 - June 2025

B.Tech - Computer Science and Engineering | CGPA - 8.45

Pune, India

Maharashtra Vidyalaya Barshi

May 2019 - June 2021

Higher Secondary Education (Science) | Percentage - 86.85%

Barshi, India

EXPERIENCE

Euron

August 2025

AI Developer Intern

Remote, India — Individual project-based

- Designed and implemented AI applications by developing models for a medical chatbot using domain-specific **data**.
- Implemented **vector search** for fast and relevant medical information retrieval.
- Developed backend services and application **systems**, containerized using Docker to support **scalable deployment**.
- Deployed on **AWS EC2** using images from AWS ECR with CI/CD automation.

CodersCave

May 2024 - June 2024

Data Science Intern

Remote, India — Individual project-based

- Developed and trained **machine learning models** on text data, achieving 80.7% accuracy.
- Streamlined text preprocessing, tokenization, and feature extraction for training pipeline.
- Evaluated and optimized **model performance**, delivering 0.72 F1-score.

PROJECTS

AutoML Brain – ML Governance & Canary Deployment Platform | [GitHub](#) Jan 2026 – Apr 2026

- Built a **3-agent autonomous ML system** (Governance Agent, Deployment Agent, GPT-4o-mini Reporter) on AWS EKS with a 7-stage DVC pipeline to automate the full ML lifecycle.
- Reduced FastAPI request-to-response P95 by 82% (2000ms → 350ms) by controlling LightGBM thread usage in K8s pods and removing preprocessing from the serving path.
- Implemented **canary deployment** with shadow traffic and Prometheus-based validation, enabling safe model promotion using latency-based thresholds.
- Developed a **GPT-4o-mini audit layer** generating PDF governance reports from model metrics, SHAP feature drivers, and PSI drift scores after every deployment.

CodeSentinel AI – Multi-Agent GenAI Code Review Platform | [GitHub](#) March 2026 – April 2026

- Architected a 7-service GenAI microservices platform on AWS EKS using LangGraph pipeline to process GitHub PR webhooks; achieved P95 webhook latency 620 ms and **0% failure** under Locust load testing.
- Built a DVC-based RAG pipeline with AST-based chunking, MiniLM + BM25 hybrid retrieval, RRF fusion, and CrossEncoder reranking over 17K+ functions; achieved **RAGAS faithfulness 0.95** and context precision 0.92.
- Optimized inference cost to **0.0037** dollar per review (93%) reduction using token-aware routing and selective model usage, with full LLM observability.
- Deployed 7 FastAPI microservices on Kubernetes EKS with autoscaling, integrated Semgrep security scanning, and implemented monitoring via Prometheus and Grafana with CI/CD pipelines.

FraudGuard – Real-Time Fraud Detection & MLOps Platform | [GitHub](#) Nov 2025 – Jan 2026

- Mitigated **0.75% class imbalance** using LightGBM with cost-sensitive learning and F1-optimized threshold via precision-recall curve, achieving PR-AUC 0.945.
- Deployed FastAPI on EKS with async Kafka producers and custom Prometheus instrumentation, sustaining **1.2K RPS** with **model-only inference P95 < 25ms and P99 < 32ms**.
- Architected a **DVC pipeline with PSI drift detection** that auto-dispatches GitHub Actions retraining with an MLflow quality gate before every Production promotion.
- Implemented NGINX canary ingress at 10% traffic with automated promote-or-rollback validation ensuring **zero-downtime releases on AWS EKS**.

TECHNICAL SKILLS

Languages: Python, SQL

Technologies: Machine Learning, Deep Learning, NLP, RAG, LangChain, AI Agents

Frameworks & Libraries: PyTorch, Langchain, SHAP

Database: SQL, Vector Database

MLOps-Tools: Git, GitHub, DVC, MLflow, Github Actions, DagsHub, Prometheus, Grafana

Deployment & Cloud Services: FastAPI, Docker, AWS, Kubernetes (EKS)