

AKASH AGALAVE

Pune, India

☎ +91-9359372603

✉ akash.agalave.work@gmail.com

🌐 [linkedin.com/in/akash-agalave](https://www.linkedin.com/in/akash-agalave)

🐙 github.com/akashagalave

EDUCATION

Dr. D.Y. Patil University

B.Tech - Computer Science and Engineering | CGPA - 8.45

Dec 2021 - June 2025

Pune, India

Maharashtra Vidyalaya Barshi

Higher Secondary Education (Science) | Percentage - 86.85%

May 2019 - June 2021

Barshi, India

EXPERIENCE

Euron

August 2025

AI Developer Intern

Remote, India — Individual project-based

- Designed and implemented AI applications by developing models for a medical chatbot using domain-specific **data**.
- Implemented **vector search** for fast and relevant medical information retrieval.
- Developed backend services and application **systems**, containerized using Docker to support **scalable deployment**.
- Deployed on **AWS EC2** using images from AWS ECR with CI/CD automation.

CodersCave

May 2024 - June 2024

Data Science Intern

Remote, India — Individual project-based

- Developed and trained **machine learning models** on text data, achieving 80.7% accuracy.
- Streamlined text preprocessing, tokenization, and feature extraction for training pipeline.
- Evaluated and optimized **model performance**, delivering 0.72 F1-score.

PROJECTS

FeedbackLens AI – Multi-Agent RAG Feedback Intelligence System | [GitHub](#)

April 2026

- Built a **multi-agent** RAG system using LangGraph to automate large-scale customer feedback analysis, replacing slow and inconsistent manual review workflows..
- Implemented a 3-agent pipeline with Qdrant and **BM25** hybrid retrieval and LLM reasoning, achieving 0.75 evaluation score using a custom deterministic framework.
- Improved retrieval precision and efficiency using query normalization, semantic filtering, issue standardization, and Redis caching, maintaining **5s** average end-to-end latency.
- Deployed** FastAPI microservices on AWS EKS with HPA autoscaling, integrating Qdrant vector database, Redis caching, and CI/CD pipelines for data ingestion and deployment.

AutoML Brain – ML Governance & Canary Deployment Platform | [GitHub](#)

Jan 2026 – Apr 2026

- Built a **3-agent autonomous ML system** (Governance Agent, Deployment Agent, GPT-4o-mini Reporter) on AWS EKS with a 7-stage DVC pipeline to automate the full ML lifecycle.
- Reduced FastAPI request-to-response P95 by 82% (2000ms → 350ms) by controlling LightGBM thread usage in K8s pods and removing preprocessing from the serving path.
- Implemented **canary deployment** with shadow traffic and Prometheus-based validation, enabling safe model promotion using latency-based thresholds.
- Developed a **GPT-4o-mini audit layer** generating PDF governance reports from model metrics, SHAP feature drivers, and PSI drift scores after every deployment.

FraudGuard – Real-Time Fraud Detection & MLOps Platform | [GitHub](#)

Nov 2025 – Jan 2026

- Mitigated **0.75% class imbalance** using LightGBM with cost-sensitive learning and F1-optimized threshold via precision-recall curve, achieving **PR-AUC 0.945**.
- Deployed FastAPI on EKS with async Kafka producers and custom Prometheus instrumentation, sustaining **1.2K RPS** with model-only inference P95 24ms and P99 31ms..
- Architected a DVC pipeline with PSI drift detection that auto-dispatches GitHub Actions retraining with an MLflow quality gate before every Production promotion.
- Implemented NGINX canary ingress at 10% traffic with automated promote-or-rollback validation ensuring **zero-downtime** releases on AWS EKS.

TECHNICAL SKILLS

Languages: Python, SQL

Technologies: Machine Learning, Deep Learning, NLP, RAG, AI Agents

Frameworks & Libraries: PyTorch, Langchain, SHAP

Database: SQL, Vector Database

MLOps-Tools: Git, GitHub, DVC, MLflow, Github Actions, DagsHub, Prometheus, Grafana

Deployment & Cloud Services: FastAPI, Docker, AWS, Kubernetes (EKS)